



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Message Passing Interface (MPI)

Summer School 2017 – Effective High Performance Computing

Tim Robinson, CSCS

July 19–20, 2017

Previous course summary

- Point-to-point communication
- Blocking and non-blocking communication
- Transfer modes

Course Objectives

- The understanding of a collective operations
- Knowledge of the different collective operations

General Course Structure



- An introduction to MPI
- Point-to-point communications
- Collective communications
- Topology
- Datatypes

General Course Structure



- An introduction to MPI
- Point-to-point communications
- Collective communications
 - Collective communications
 - Barrier
 - Broadcast
 - Scatter/Gather
 - All to all
 - Reduction
 - Global collective operations
 - Non-blocking coll-op
- Topology
- Datatypes



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich

Collective communications

Collective operations

Communications involving a group of processes part of a communicator.
Different algorithms: $1 \rightarrow N$, $N \rightarrow 1$ or $N \rightarrow N$ ($1 \rightarrow 1 = \text{pt2pt}$).

Example:

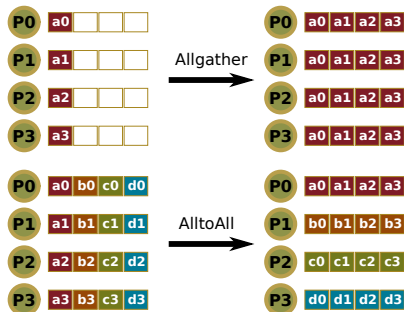

- Barrier Synchronization
- Broadcast
- Gather/Scatter
- AlltoAll
- Reduction (sum, max, prod, ...)

Features:

- All processes must call the collective routine, one is the root
- No tags

The MPI library should use the most efficient communication algorithm for the particular platform.

Collective operations schemes




Barrier

Stop processes until all processes within a communicator reach the barrier.

Pseudo-code

```
MPI_Barrier(comm)
```



Broadcast

One-to-all communication: same data sent from root process to all other processes in the communicator.

Pseudo-code

```
MPI_Bcast(buf, count, type, root, comm)
```

root rank being the initiator of the collective operation

Scatter

One-to-all communication: different data sent from the root process to all other processes in the communicator.

Pseudo-code

```
MPI_Scatter(sndbuf, sndcount, sndtype,  
           rcvbuf, rcvcount, rcvtype, root, comm)
```

| | |
|-----------------|--|
| sndcount | number of elements sent to each process, not the size of <code>sndbuf</code> , that should be <code>sndcount</code> times the number of process in the communicator |
| rcvcount | number of element in the receive buffer |

The sender arguments are meaningful only for root.

Gather

One-to-all communication: different data collected by the root process, from all others processes in the communicator.

Pseudo-code

```
MPI_Gather(sndbuf, sndcount, sndtype,  
          rcvbuf, rcvcount, rcvtype, root, comm)
```

rcvcount the number of elements collected from each process,
not the size of rcvbuf, that should be rcvcount times
the number of process in the communicator

sndcount number of element in the send buffer

The receive arguments are meaningful only for root.

Global exchange: All to All

All-to-all communication: global exchange, all processes exchange their data. Useful for data transposition.

Pseudo-code

```
MPI_Alltoall(sndbuf, sndcount, sndtype,  
             rcvbuf, rcvcount, rcvtype, comm)
```

Reduction

The reduction operation allows to:

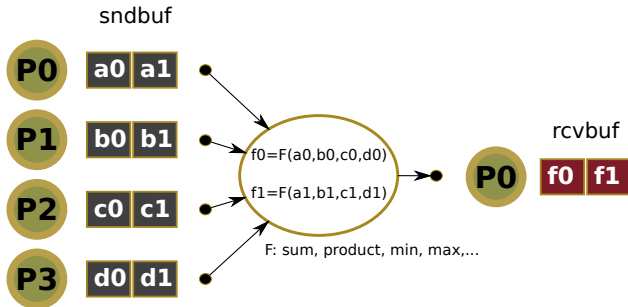
- Collect data from each process
- Reduce the data to a single value
- Store the result on the root processes
- Store the result on all processes
- Overlap communication and computation

Reduction

Pseudo-code

```
MPI_Reduce(sndbuf, rcvbuf, count, type, op, root, comm)
```

op parallel operation to perform



Reduction operators

| MPI op | Operation |
|------------|----------------------|
| MPI_MAX | Maximum |
| MPI_MIN | Minimum |
| MPI_SUM | Sum |
| MPI_PROD | Product |
| MPI_LAND | Logical AND |
| MPI_BAND | Bitwise AND |
| MPI_LOR | Logical OR |
| MPI_BOR | Bitwise OR |
| MPI_LXOR | Logical exclusive OR |
| MPI_BXOR | Bitwise exclusive OR |
| MPI_MAXLOC | Maximum and location |
| MPI_MINLOC | Minimum and location |

Global collective operations

The result of the one-to-all operation is known by all ranks at the end of the operation.

Pseudo-code

```
MPI_Allgather(sndbuf, sndcount, sndtype,  
              rcvbuf, rcvcount, rcvtype, comm)  
  
MPI_Allreduce(sndbuf, rcvbuf, count, type, op, comm)
```

The argument **root** is missing, the result is stored in all processes.

Non-blocking collective operations

All collective operations have a non-blocking version.

Example:

Pseudo-code

```
MPI_Ibcast(buf, count, type, root, comm, request)
```

Other functions:

Pseudo-code

```
MPI_Ibarrier, MPI_Igather, MPI_Ireduce, MPI_Iscatter,  
MPI_Iallgather, MPI_Iallreduce, MPI_Ialltoall
```

Other functions

- Operations with different buffer sizes:

```
MPI_AlltoAllv, MPI_Gatherv, MPI_Scatterv, MPI_Allgatherv
```

- Neighbor operations, based on topology:

```
MPI_Neighbor_gather, MPI_Neighbor_alltoall
```

- Cumulative per rank reduction:

```
MPI_Scan, MPI_Exscan
```

- Create your own operator:

```
MPI_Op_create, MPI_Op_free
```

Practicals

Exercise: 03.MPI_Coll

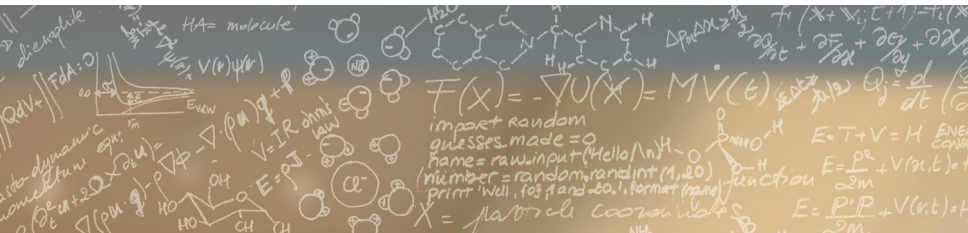
1. Read from the terminal and broadcast the input
2. Initialise an array and scatter it
3. Reduction operation
4. Reduction with results stored in all ranks (allreduce)



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Thank you for your attention.