

# Introduction to HDF5

Adapted from presentation by  
the HDF5 Group

1

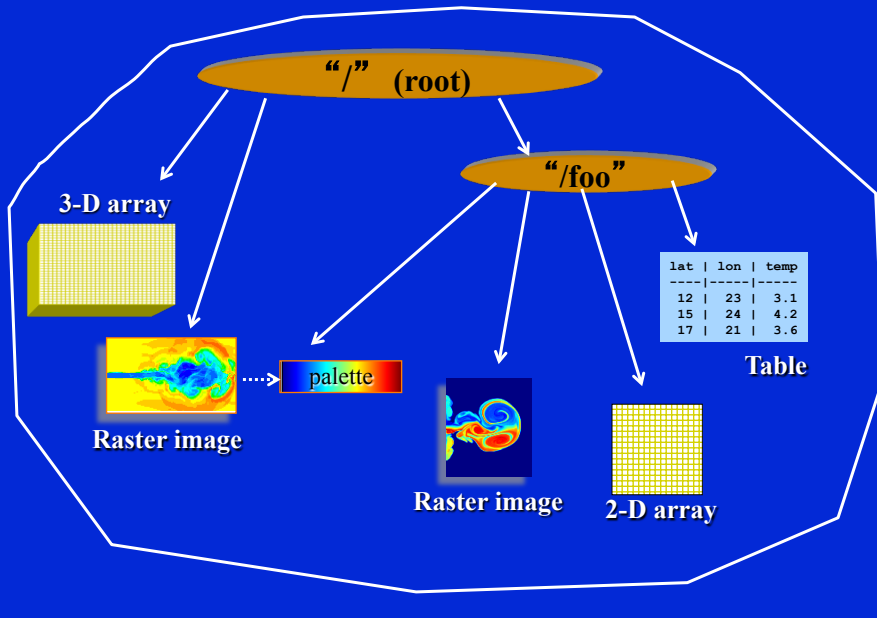
## What is HDF5?

---

- **File format for storing (scientific) data**
  - To store and organize all kinds of (scientific) data
  - To share data, to port files from one platform to another
  - To overcome a limit on number and size of the objects in the file
- **Software for accessing (scientific) data**
  - Flexible I/O library (parallel, remote, etc.)
  - Efficient storage
  - Available on almost all platforms
  - C, F90, C++ , Java, Python (h5py) APIs
  - Tools (HDFView, utilities)

2

## Example HDF5 file



## HDF5 file

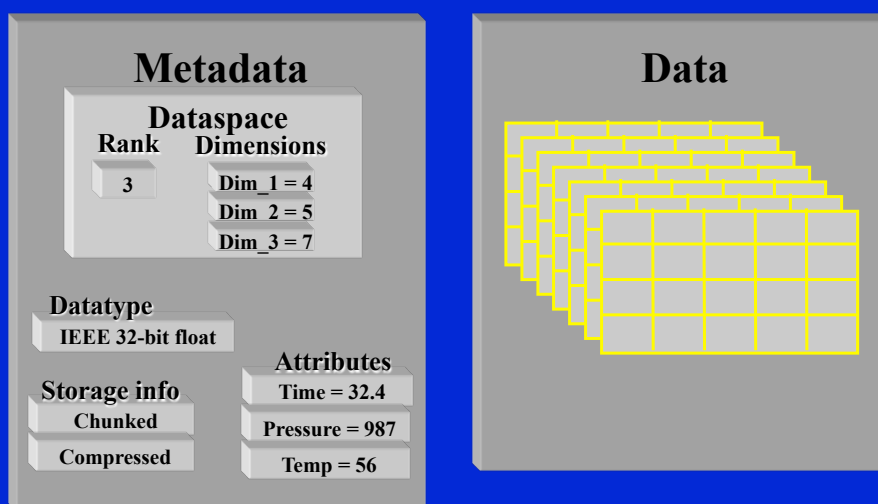
- **Primary Objects**
  - Groups
  - Datasets
- **Additional means to organize data**
  - Attributes
  - Sharable objects
  - Storage and access properties

## HDF5 Dataset

- **Data array**
  - ordered collection of identically typed data items distinguished by their indices
- **Metadata**
  - Dataspace – rank, dimensions, other spatial info about dataset
  - Datatype
  - Attribute list – user-defined metadata
  - Special storage options – how array is organized


6

## Dataset Components



## Dataspaces


- **Dataspace – spatial info about a dataset**
  - Rank and dimensions
    - Permanent part of dataset definition
  - Subset of points, for partial I/O
    - Needed only during I/O operations
- **Apply to datasets in memory or in the file**




Rank = 2  
Dimensions = 4x6

8


## Sample Mappings between File Dataspaces and Memory Dataspaces




(a) Hyperslab from a 2D array to the corner of a smaller 2D array



(b) Regular series of blocks from a 2D array to a contiguous sequence at a certain offset in a 1D array



(c) A sequence of points from a 2D array to a sequence of points in a 3D array.



(d) Union of hyperslabs in file to union of hyperslabs in memory.

9

## **Datatypes (array elements)**

---

- **Datatype – how to interpret a data element**
  - Permanent part of the dataset definition
- **HDF5 atomic types**
  - normal integer & float
  - user-definable integer and float (e.g. 13-bit integer)
  - variable length types (e.g. strings)
  - pointers - references to objects/dataset regions
  - enumeration - names mapped to integers
  - array
- **HDF5 compound types**
  - Comparable to C structs
  - Members can be atomic or compound types

10

## **Attributes**


---

- **Attribute – data of the form “name = value”, attached to an object**
- **Operations are scaled-down versions of the dataset operations**
  - Not extendible
  - No compression
  - No partial I/O
- **Optional for the dataset definition**
- **Can be overwritten, deleted, added during the “life” of a dataset**

11


## Special Storage Options

**chunked**




Better subsetting  
access time;  
extendable

**compressed**




Improves storage  
efficiency,  
transmission speed

**extendable**



Arrays can be  
extended in any  
direction


**External  
file**



Metadata in one file,  
raw data in another.


## Groups

- **Group – a mechanism for describing collections of related objects**
- Every file starts with a root group
- Can have attributes
- Similar to UNIX directories, but cycles are allowed




## HDF5 objects and their pathnames

/ (root)  
/x  
/foo  
/foo/temp  
/foo/bar/temp



14

## Shared groups & members




15

## HDF5 I/O Library

16

### Structure of HDF5 Library

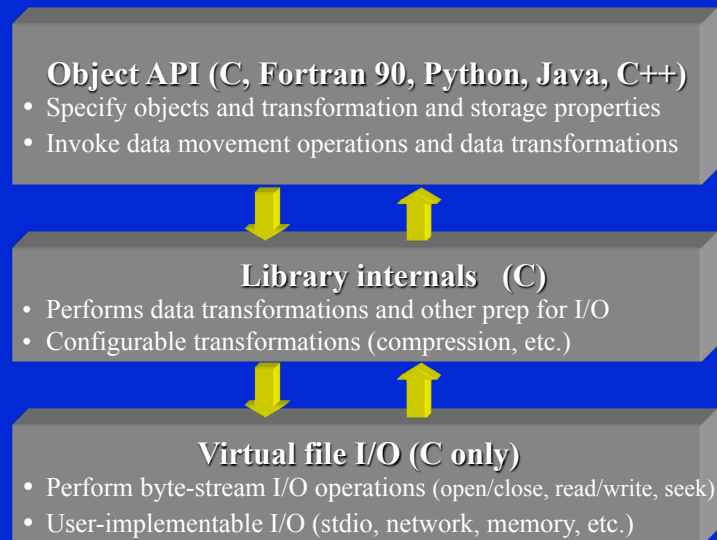
---



17



## Structure of HDF5 Library



18

## The General HDF5 API

- **Currently has C, Fortran 90, Python, Java and C++ bindings.**
- **C routines begin with prefix H5\*, where \* is a single letter indicating the object on which the operation is to be performed.**
- **Full functionality**
- **Fortran: same name followed by “\_f”**

Example APIs:

**H5D** : Dataset interface    e.g.. **H5Dread**  
**H5F** : File interface        e.g.. **H5Fopen**  
**H5S** : dataSpace interface e.g.. **H5Sclose**

19

## The General Paradigm

---

- *Properties (called creation and access property lists) of objects are defined*
- **Objects are opened or created**
- **Objects then accessed**
- **Objects finally closed**

20

## HDF5 C Programming Issues

---

For portability, HDF5 library has its own defined types:

<b>hid_t:</b>	object identifiers (native <i>integer</i> )
<b>hsize_t:</b>	size used for dimensions ( <i>unsigned long</i> or <i>unsigned long long</i> )
<b>hssize_t:</b>	for specifying coordinates and sometimes for dimensions ( <i>signed long</i> or <i>signed long long</i> )
<b>herr_t:</b>	function return value
<b>hvl_t:</b>	variable length datatype

For **C**, include `#include hdf5.h` at the top of your HDF5 application. For F90: USE HDF5.

21

---

# Files

24

## Example 1

```
1  hid_t      file_id;  
2  herr_t      status;  
3  file_id = H5Fcreate ("file.h5", H5F_ACC_TRUNC,  
                      H5P_DEFAULT, H5P_DEFAULT);  
4  status = H5Fclose (file_id);
```

Create a new file using  
default properties



### Example 1

```
1  hid_t      file_id;
2  herr_t      status;

3  file_id = H5Fcreate ("file.h5", H5F_ACC_TRUNC,
                      H5P_DEFAULT, H5P_DEFAULT);

4  status = H5Fclose (file_id);
```



Terminate access to  
the File

26

### h5\_crfile.c

```
1  #include <hdf5.h>
2  #define FILE "file.h5"
3
4  main() {
5
6      hid_t      file_id;  /* file identifier */
7      herr_t      status;
8
9      /* Create a new file using default properties. */
10     file_id = H5Fcreate (FILE, H5F_ACC_TRUNC,
                          H5P_DEFAULT, H5P_DEFAULT);
11
12     /* Terminate access to the file. */
13     status = H5Fclose (file_id);
14 }
```

## Example 1: h5dump Output

---

```
HDF5 "file.h5" {  
  GROUP "/" {  
  }  
}
```



'/'

28

---

## Groups

29

## Steps to use groups

---

An HDF5 group is a structure containing zero or more HDF5 objects. The two primary HDF5 objects are groups and datasets. To create a group, the calling program must:

- Obtain the location identifier where the group is to be created.
- Create the group.
- Close the group.

30

---

```
hid_t H5Gcreate( hid_t loc_id, const char *name,  
                size_t size_hint )
```

**loc\_id:** File or parent group identifier.

**name:** Absolute or relative name of the new group.

**hint:** number of bytes to reserve for the names that will appear in the group (0 ok)

31

## Group usage

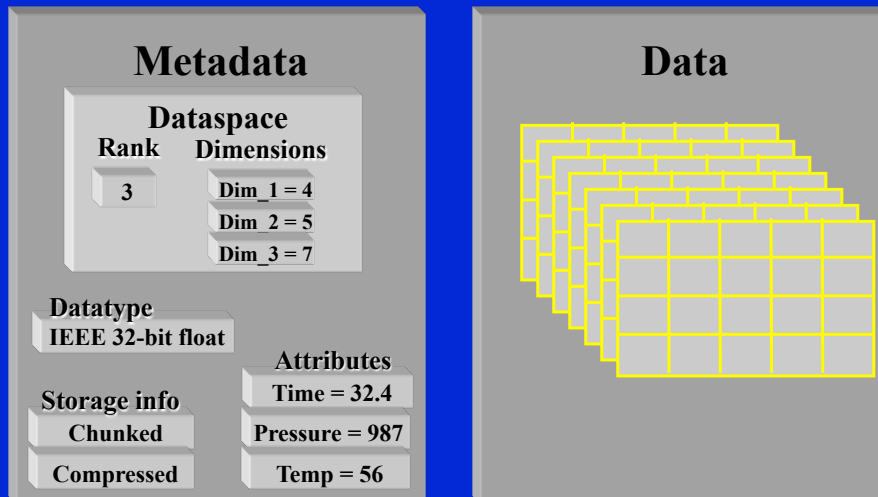
---

```
CALL h5open_f(error)
!  
! Create a new file using default properties.  
!  
CALL h5fcreate_f(filename, H5F_ACC_TRUNC_F, file_id, error)  
  
!  
! Create a group named "/MyGroup" in the file.  
!  
CALL h5gcreate_f(file_id, groupname, group_id, error)  
  
!  
! Close the group.  
!  
CALL h5gclose_f(group_id, error)  
  
!  
! Terminate access to the file.  
!  
CALL h5fclose_f(file_id, error)  
  
Close FORTRAN interface.  
  
CALL h5close_f(error)
```

32

## Datasets

## Dataset Components



34

### Example 2 – Create an empty 4x6 dataset

```

1 hid_t      file_id, dataset_id, dataspace_id;
2 hsize_t    dims[2];
3 herr_t     status;
4 file_id = H5Fcreate ("dset.h5", H5F_ACC_TRUNC,
                      H5P_DEFAULT, H5P_DEFAULT);

5 dims[0] = 4;
6 dims[1] = 6;
7 dataspace_id = H5Screate_simple (2, dims, NULL);

8 dataset_id = H5Dcreate(file_id, "dset", H5T_STD_I32BE,
                      dataspace_id, H5P_DEFAULT);

9 status = H5Dclose (dataset_id);
10 status = H5Sclose (dataspace_id);
11 status = H5Fclose (file_id);

```



### Example 2 – Create an empty 4x6 dataset

```
1 hid_t      file_id, dataset_id, dataspace_id;
2 hsize_t    dims[2];
3 herr_t      status;

4 file_id = H5Fcreate ("dset.h5", H5F_ACC_TRUNC,
  Create a dataspace      H5P_DEFAULT, H5P_DEFAULT);
5 dims[0] = 4;
6 dims[1] = 6;
7 dataspace_id = H5Screate_simple (2, dims, NULL);

8 dataset_id = H5Dcreate(file_id,"dset",H5T_STD_I32BE,
                        dataspace_id, H5P_DEFAULT);

9 status = H5Dclose (dataset_id);
10 status = H5Sclose (dataspace_id);
11 status = H5Fclose (file_id);
```

### Example 2 – Create an empty 4x6 dataset

```
1 hid_t      file_id, dataset_id, dataspace_id;
2 hsize_t    dims[2];
3 herr_t      status;

4 file_id = H5Fcreate ("dset.h5", H5F_ACC_TRUNC,
  Create a dataspace      H5P_DEFAULT, H5P_DEFAULT);
5 dims[0] = 4;
6 dims[1] = 6;
7 dataspace_id = H5Screate_simple (2, dims, NULL);
8 dataset_id = H5Dcreate(file_id,"dset",H5T_STD_I32BE,
                        dataspace_id, H5P_DEFAULT);

9 status = H5Dclose (dataset_id);
10 status = H5Sclose (dataspace_id);
11 status = H5Fclose (file_id);
```

rank

current dims

Set maxdims to current dims

### Example 2 – Create an empty 4x6 dataset

```
1 hid_t      file_id, dataset_id, dataspace_id;
2 hsize_t    dims[2];
3 herr_t     status;

4 file_id = H5Fcreate ("dset.h5", H5F_ACC_TRUNC,
                      H5P_DEFAULT, H5P_DEFAULT);

5 dims[0] = 4;
6 dims[1] = 6;
7 dataspace_id = H5Screate_simple (2, dims, NULL);

Create a dataset

8 dataset_id = H5Dcreate(file_id,"dset",H5T_STD_I32BE,
                        dataspace_id, H5P_DEFAULT);

9 status = H5Dclose (dataset_id);
10 status = H5Sclose (dataspace_id);
11 status = H5Fclose (file_id);
```

### Example 2 – Create an empty 4x6 dataset

```
1 hid_t      file_id, dataset_id, dataspace_id;
2 hsize_t    dims[2];
3 herr_t     status;

4 file_id = H5Fcreate ("dset.h5", H5F_ACC_TRUNC,
                      H5P_DEFAULT, H5P_DEFAULT);

5 dims[0] = 4;
6 dims[1] = 6;
7 dataspace_id = H5Screate_simple (2, dims, NULL);

Create a dataset

8 dataset_id = H5Dcreate(file_id,"dset",H5T_STD_I32BE,
                        dataspace_id, H5P_DEFAULT);

9 status = H5Dclose (dataset_id);
10 status = H5Sclose (dataspace_id);
11 status = H5Fclose (file_id);
```

Pathname

Datatype

Dataspace

Property list (default)

### Example 2 – Create an empty 4x6 dataset

```
1 hid_t      file_id, dataset_id, dataspace_id;
2 hsize_t    dims[2];
3 herr_t      status;

4 file_id = H5Fcreate ("dset.h5", H5F_ACC_TRUNC,
                      H5P_DEFAULT, H5P_DEFAULT);

5 dims[0] = 4;
6 dims[1] = 6;
7 dataspace_id = H5Screate_simple (2, dims, NULL);


8 dataset_id = H5Dcreate(file_id,"dset",H5T_STD_I32BE,
                        dataspace_id, H5P_DEFAULT);
Terminate access to dataset, dataspace, & file

9 status = H5Dclose (dataset_id);
10 status = H5Sclose (dataspace_id);
11 status = H5Fclose (file_id);
```

### Example2: h5dump Output

An empty 4x6 dataset

```
HDF5 "dset.h5" {
GROUP "/" {
  DATASET "dset" {
    DATATYPE { H5T_STD_I32BE }
    DATASPACE { SIMPLE ( 4, 6 ) / ( 4, 6 ) }
    DATA {
      0, 0, 0, 0, 0, 0,
      0, 0, 0, 0, 0, 0,
      0, 0, 0, 0, 0, 0,
      0, 0, 0, 0, 0, 0
    }
  }
}
}
```



# Writing and Reading Datasets

## Dataset I/O

---

- **Dataset I/O involves**
  - reading or writing
  - all or part of a dataset
  - Compressed/uncompressed
- **During I/O operations data is translated between the source & destination (file-memory, memory-file)**
  - Datatype conversion
    - data types (e.g. 16-bit integer => 32-bit integer) of the same class
  - Dataspace conversion
    - dataspace (e.g. 10x20 2d array => 200 1d array)

55

## Partial I/O

---

- **Selected elements (called selections) from source are mapped (read/written) to the selected elements in destination**
- **Selection**
  - Selections in memory can differ from selection in file
  - Number of selected elements is always **the same** in source and destination
- **Selection can be**
  - Hyperslabs (contiguous blocks, regularly spaced blocks)
  - Points
  - Results of set operations (union, difference, etc.) on hyperslabs or points

56

## Reading Dataset into Memory from File

---

**File**

*2D array of 16-bit ints*




**Memory**

*3D array of 32-bit ints*




57

## Reading Dataset into Memory from File




58

## Reading Dataset into Memory from File



59

## Reading Dataset into Memory from File



60

## Steps for Dataset Writing/Reading

1. If necessary, open the file to obtain the file ID
2. Open the dataset to obtain the dataset ID
3. Specify
  - Memory datatype
  - Library “*knows*” file *datatype* – do not need to specify
  - Memory dataspace
  - File dataspace
  - Transfer properties (optional)
4. Perform the desired operation on the dataset
5. Close dataspace, datatype and property lists

61

### Example 3 – Writing to an existing dataset

```
1 hid_t      file_id, dataset_id;
2 herr_t     status;
3 int        i, j, dset_data[4][6];

4 for (i = 0; i < 4; i++)
5     for (j = 0; j < 6; j++)
6         dset_data[i][j] = i * 6 + j + 1;

7 file_id = H5Fopen ("dset.h5", H5F_ACC_RDWR, H5P_DEFAULT);
8 dataset_id = H5Dopen (file_id, "dset");

9 status = H5Dwrite (dataset_id, H5T_NATIVE_INT,
                    H5S_ALL, H5S_ALL, H5P_DEFAULT, dset_data);
```

### Example 3 – Writing to an existing dataset

```
1 hid_t      file_id, dataset_id;
2 herr_t     status;
3 int        i, j, dset_data[4][6];

4 Initialize buffer

4 for (i = 0; i < 4; i++)
5     for (j = 0; j < 6; j++)
6         dset_data[i][j] = i * 6 + j + 1;

7 file_id = H5Fopen ("dset.h5", H5F_ACC_RDWR, H5P_DEFAULT);
8 dataset_id = H5Dopen (file_id, "dset");

9 status = H5Dwrite (dataset_id, H5T_NATIVE_INT,
                    H5S_ALL, H5S_ALL, H5P_DEFAULT, dset_data);
```



### Example 3 – Writing to an existing dataset

```
1 hid_t      file_id, dataset_id;
2 herr_t      status;
3 int         i, j, dset_data[4][6];

4 for (i = 0; i < 4; i++)
5     for (j = 0; j < 6; j++)
6         dset_data[i][j] = i * 6 + j + 1;

7 file_id = H5Fopen ("dset.h5", H5F_ACC_RDWR, H5P_DEFAULT);
8 dataset_id = H5Dopen (file_id, "dset");

9 status = H5Dwrite (dataset_id, H5T_NATIVE_INT,
                    H5S_ALL, H5S_ALL, H5P_DEFAULT, dset_data);
```

#### Open existing file and dataset

### Example 3 – Writing to an existing dataset

```
1 hid_t      file_id, dataset_id;
2 herr_t      status;
3 int         i, j, dset_data[4][6];

4 for (i = 0; i < 4; i++)
5     for (j = 0; j < 6; j++)
6         dset_data[i][j] = i * 6 + j + 1;

7 file_id = H5Fopen ("dset.h5", H5F_ACC_RDWR, H5P_DEFAULT);
8 dataset_id = H5Dopen (file_id, "dset");

9 status = H5Dwrite (dataset_id, H5T_NATIVE_INT,
                    H5S_ALL, H5S_ALL, H5P_DEFAULT, dset_data);
```

#### Write to dataset

### Example 3: h5dump Output

---

```
HDF5 "dset.h5" {  
  GROUP "/" {  
    DATASET "dset" {  
      DATATYPE { H5T_STD_I32BE }  
      DATASPACE { SIMPLE ( 4, 6 ) / ( 4, 6 ) }  
      DATA {  
        1, 2, 3, 4, 5, 6,  
        7, 8, 9, 10, 11, 12,  
        13, 14, 15, 16, 17, 18,  
        19, 20, 21, 22, 23, 24  
      }  
    }  
  }  
}
```

69